



Full Training Dynamics (20k iterations)


Training Loss




Attention Pattern Emergence




All Heads → Position 0



Gradient by Position (snapshots)



Head-Specific Gradients Over Training



Head 7: Attention vs Gradient

